

agentic-memory

Long-Context Benchmark Report

Benchmark evidence for >300k-token conversations comparing short-context prompting, full transcript stuffing, and agentic-memory retrieval.

Pass Rate

58.3%

memory_enabled

Prompt Tokens

289.2

avg per answer

Cost vs Full

152.2x cheaper

memory_enabled

Headline

On >300k-token conversations, agentic-memory retrieval outperformed both short-context and full-context baselines while using dramatically fewer prompt tokens than transcript stuffing.

Key Findings

- ? memory_enabled achieved a 7/12 pass rate (58.3%), versus 3/12 (25.0%) for full_context and 0/12 (0.0%) for short_context.
- ? memory_enabled used 289.2 prompt tokens on average, compared with 162590.2 for full_context.
- ? memory_enabled was 50.1% faster on average than full_context.
- ? full_context cost about 152.2x more per example than memory_enabled.

Mode Comparison

Measured on 12 beyond-250k conversations with Gemini 2.5 Flash.

Mode	Pass	Judge	Latency	Prompt Tokens	Cost
memory_enabled	7/12 (58.3%)	0.83	3470 ms	289.2	\$0.000323
full_context	3/12 (25.0%)	0.44	6952 ms	162590.2	\$0.049154
short_context	0/12 (0.0%)	0.35	3801 ms	1647.8	\$0.000827

Business Impact

- ? Accuracy lift over full_context: 33.3 percentage points.
- ? Latency reduction versus full_context: 50.1%.
- ? Prompt-token reduction versus full_context: 99.8%.
- ? Estimated cost reduction versus full_context: 99.3%.

Interpretation

On transcripts that grow beyond a practical prompt budget, brute-force context stuffing becomes expensive and still leaves quality on the table. Retrieval through agentic-memory preserves the relevant profile facts, prior events, and successful procedures while keeping prompts compact enough to stay economical.

Recommended Next Step

Tune the memory-mode answer scaffold so Gemini is required to explicitly cover every gold constraint that appears in retrieved memory. That should push the pass rate higher without giving up the cost and latency advantage.